

# IMPORT DATASET

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, roc_auc_score, auc, accuracy_score, precision_score, recall_score, f1_
import warnings
warnings.filterwarnings('ignore')
```

# LOADING THE CSV FILE (DATASET)

```
In [2]: df = pd.read_csv('Heart_data.csv')
df.head()
```

```
Out[2]:   index  id  age  gender  height  weight  ap_hi  ap_lo  cholesterol  gluc  smoke  alco  active  cardio
0      0    0  18393     2     168    62.0    110     80          1     1      0      0      1      0
1      1    1  20228     1     156    85.0    140     90          3     1      0      0      1      1
2      2    2  18857     1     165    64.0    130     70          3     1      0      0      0      1
3      3    3  17623     2     169    82.0    150    100          1     1      0      0      1      1
4      4    4  17474     1     156    56.0    100     60          1     1      0      0      0      0
```

# DATA PREPROCESSING: DATA CLEARNING AND FEATURE ENGINEERING

1. Remove id, and index column
2. Add BMI, Substance used and new age as the new feature

3. Only medically realistic range for blood pressure and ensuring range for blood pressure and ensuring systolic  $\geq$  diastolic

```
In [3]: #Create BMI column
df = df.drop('index', axis=1)
df = df.drop(columns=['id'])
df['bmi'] = df['weight'] / (df['height']/100)**2
df['age'] = (df['age'] / 365).astype(int)
df['substance used'] = (df['alco'] + df['smoke'])
```

```
In [4]: # Blood Pressure Cleaning

df = df[(df['ap_hi'] >= 80) & (df['ap_hi'] <= 250)]
df = df[(df['ap_lo'] >= 50) & (df['ap_lo'] <= 200)]

#Ensure systolic >= diastolic
df = df[df['ap_hi'] >= df['ap_lo']]

print("Remaining rows after blood pressure cleaning:", df.shape[0])
```

Remaining rows after blood pressure cleaning: 68652

```
In [5]: df.columns
```

```
Out[5]: Index(['age', 'gender', 'height', 'weight', 'ap_hi', 'ap_lo', 'cholesterol',
   'gluc', 'smoke', 'alco', 'active', 'cardio', 'bmi', 'substance used'],
  dtype='object')
```

```
In [6]: df.shape
```

```
Out[6]: (68652, 14)
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 68652 entries, 0 to 69999
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               68652 non-null   int32  
 1   gender             68652 non-null   int64  
 2   height              68652 non-null   int64  
 3   weight              68652 non-null   float64 
 4   ap_hi              68652 non-null   int64  
 5   ap_lo              68652 non-null   int64  
 6   cholesterol         68652 non-null   int64  
 7   gluc                68652 non-null   int64  
 8   smoke               68652 non-null   int64  
 9   alco                68652 non-null   int64  
 10  active              68652 non-null   int64  
 11  cardio              68652 non-null   int64  
 12  bmi                 68652 non-null   float64 
 13  substance used     68652 non-null   int64  
dtypes: float64(2), int32(1), int64(11)
memory usage: 7.6 MB
```

In [8]: `df.head()`

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	alco	active	cardio	bmi	substance used	
0	50	2	168	62.0	110	80		1	1	0	0	1	0	21.967120	0
1	55	1	156	85.0	140	90		3	1	0	0	1	1	34.927679	0
2	51	1	165	64.0	130	70		3	1	0	0	0	1	23.507805	0
3	48	2	169	82.0	150	100		1	1	0	0	1	1	28.710479	0
4	47	1	156	56.0	100	60		1	1	0	0	0	0	23.011177	0

In [9]: `df.tail()`

Out[9]:

	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	alco	active	cardio	bmi	substance used	
<b>69995</b>	52	2	168	76.0	120	80		1	1	1	0	1	0	26.927438	1
<b>69996</b>	61	1	158	126.0	140	90		2	2	0	0	1	1	50.472681	0
<b>69997</b>	52	2	183	105.0	180	90		3	1	0	1	0	1	31.353579	1
<b>69998</b>	61	1	163	72.0	135	80		1	2	0	0	0	1	27.099251	0
<b>69999</b>	56	1	170	72.0	120	80		2	1	0	0	1	0	24.913495	0

In [10]:

df.describe().T

Out[10]:

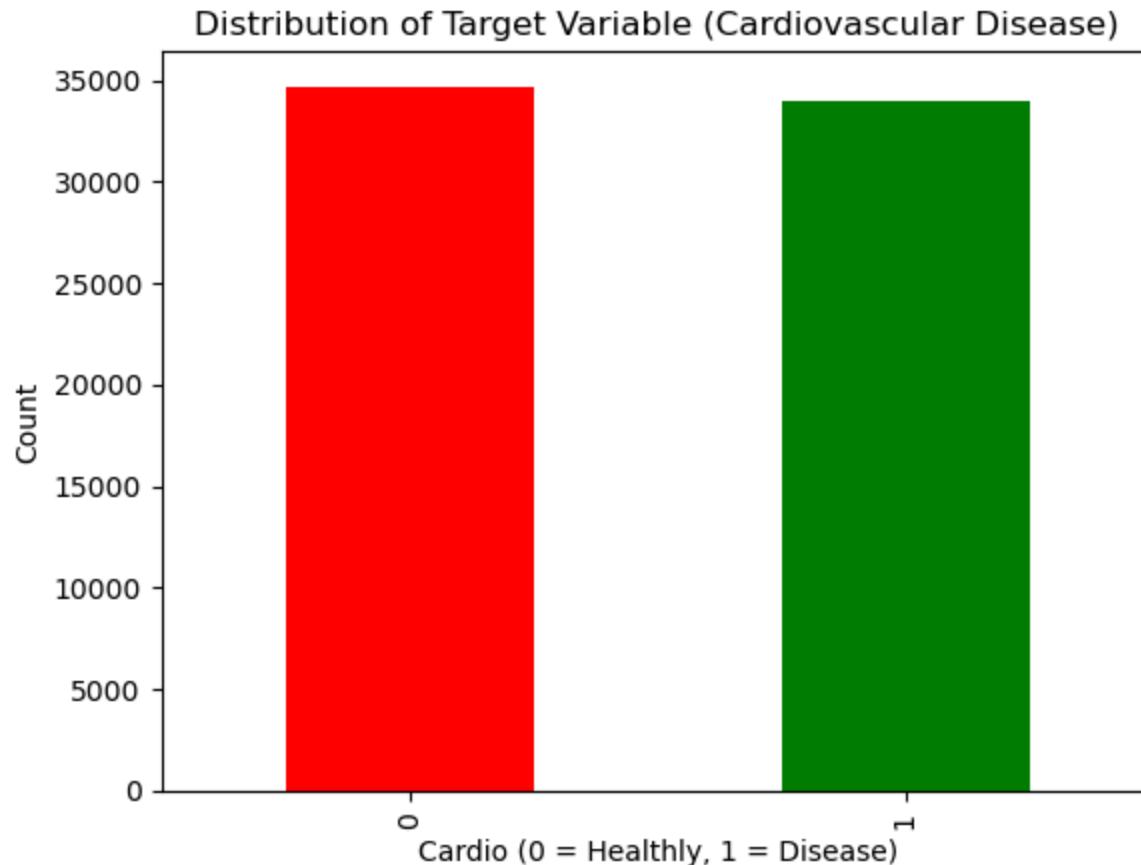
	<b>count</b>	<b>mean</b>	<b>std</b>	<b>min</b>	<b>25%</b>	<b>50%</b>	<b>75%</b>	<b>max</b>
<b>age</b>	68652.0	52.829371	6.768919	29.000000	48.000000	53.000000	58.000000	64.000000
<b>gender</b>	68652.0	1.348686	0.476558	1.000000	1.000000	1.000000	2.000000	2.000000
<b>height</b>	68652.0	164.362830	8.183262	55.000000	159.000000	165.000000	170.000000	250.000000
<b>weight</b>	68652.0	74.121230	14.331623	11.000000	65.000000	72.000000	82.000000	200.000000
<b>ap_hi</b>	68652.0	126.677198	16.686610	80.000000	120.000000	120.000000	140.000000	240.000000
<b>ap_lo</b>	68652.0	81.320457	9.430504	50.000000	80.000000	80.000000	90.000000	182.000000
<b>cholesterol</b>	68652.0	1.364651	0.678897	1.000000	1.000000	1.000000	2.000000	3.000000
<b>gluc</b>	68652.0	1.225689	0.571563	1.000000	1.000000	1.000000	1.000000	3.000000
<b>smoke</b>	68652.0	0.087965	0.283247	0.000000	0.000000	0.000000	0.000000	1.000000
<b>alco</b>	68652.0	0.053341	0.224715	0.000000	0.000000	0.000000	0.000000	1.000000
<b>active</b>	68652.0	0.803356	0.397464	0.000000	1.000000	1.000000	1.000000	1.000000
<b>cardio</b>	68652.0	0.494712	0.499976	0.000000	0.000000	0.000000	1.000000	1.000000
<b>bmi</b>	68652.0	27.522085	6.051969	3.471784	23.875115	26.346494	30.119376	298.666667
<b>substance used</b>	68652.0	0.141307	0.417219	0.000000	0.000000	0.000000	0.000000	2.000000

## EXPLORATORY DATA ANALYSIS (EDA)

1. Target Distribution
2. Demographics (Age, Gender)
3. Anthropometrics (Height, Weight, BMI)
4. Medical Indicators (Blood Pressure, Cholesterol, Glucose)
5. Lifestyle Factors (Smoke, Alcohol, Activity)
6. Correlation
7. Feature Importance

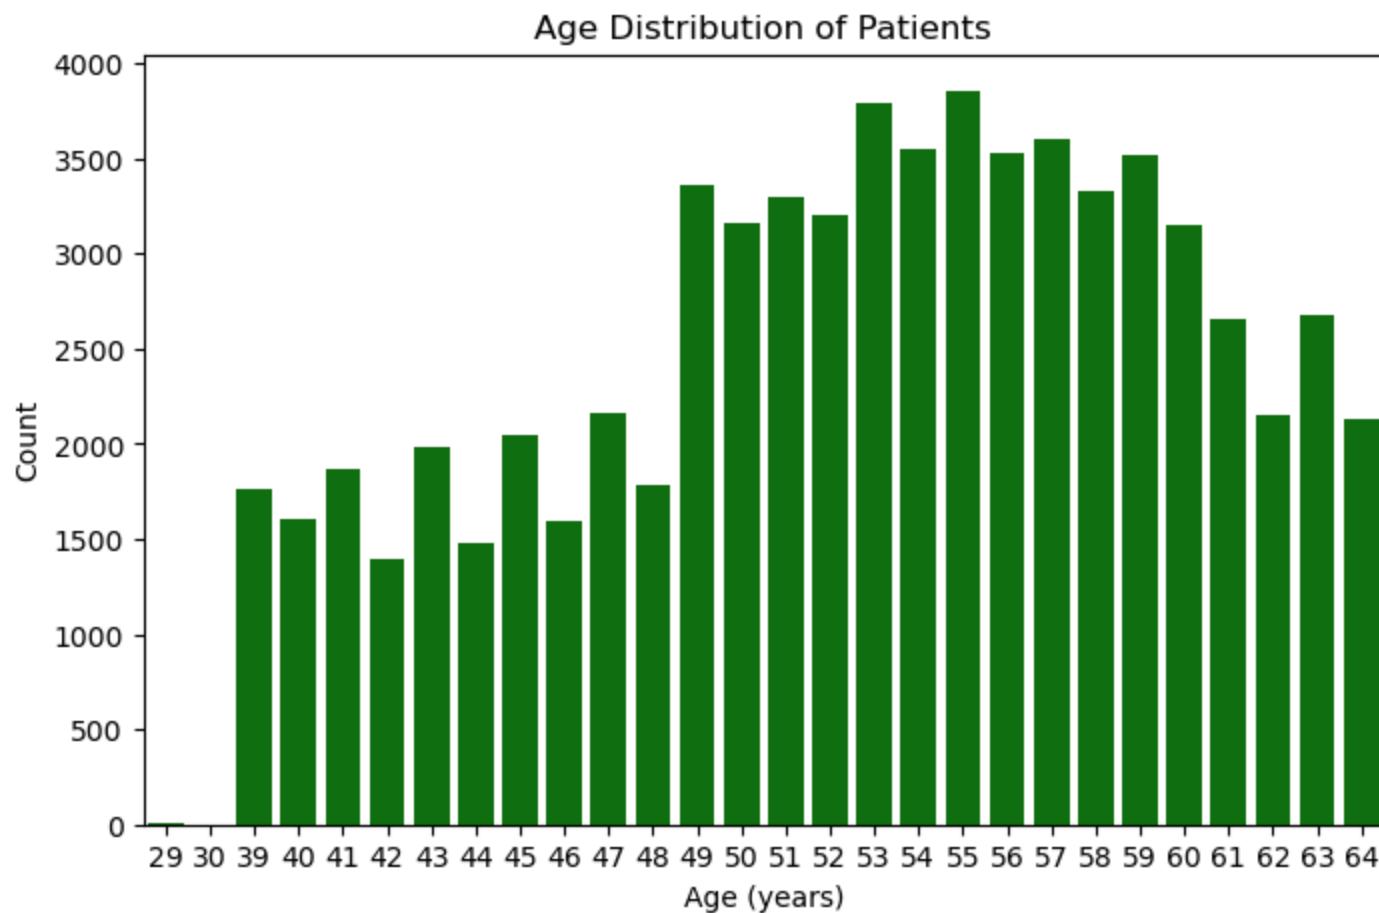
# 1. Target Variable Distribution

```
In [11]: # Countplot of target variable
df['cardio'].value_counts().plot(kind='bar', color=['red','green'])
plt.title('Distribution of Target Variable (Cardiovascular Disease)')
plt.xlabel('Cardio (0 = Healthly, 1 = Disease)')
plt.ylabel('Count')
plt.show()
```



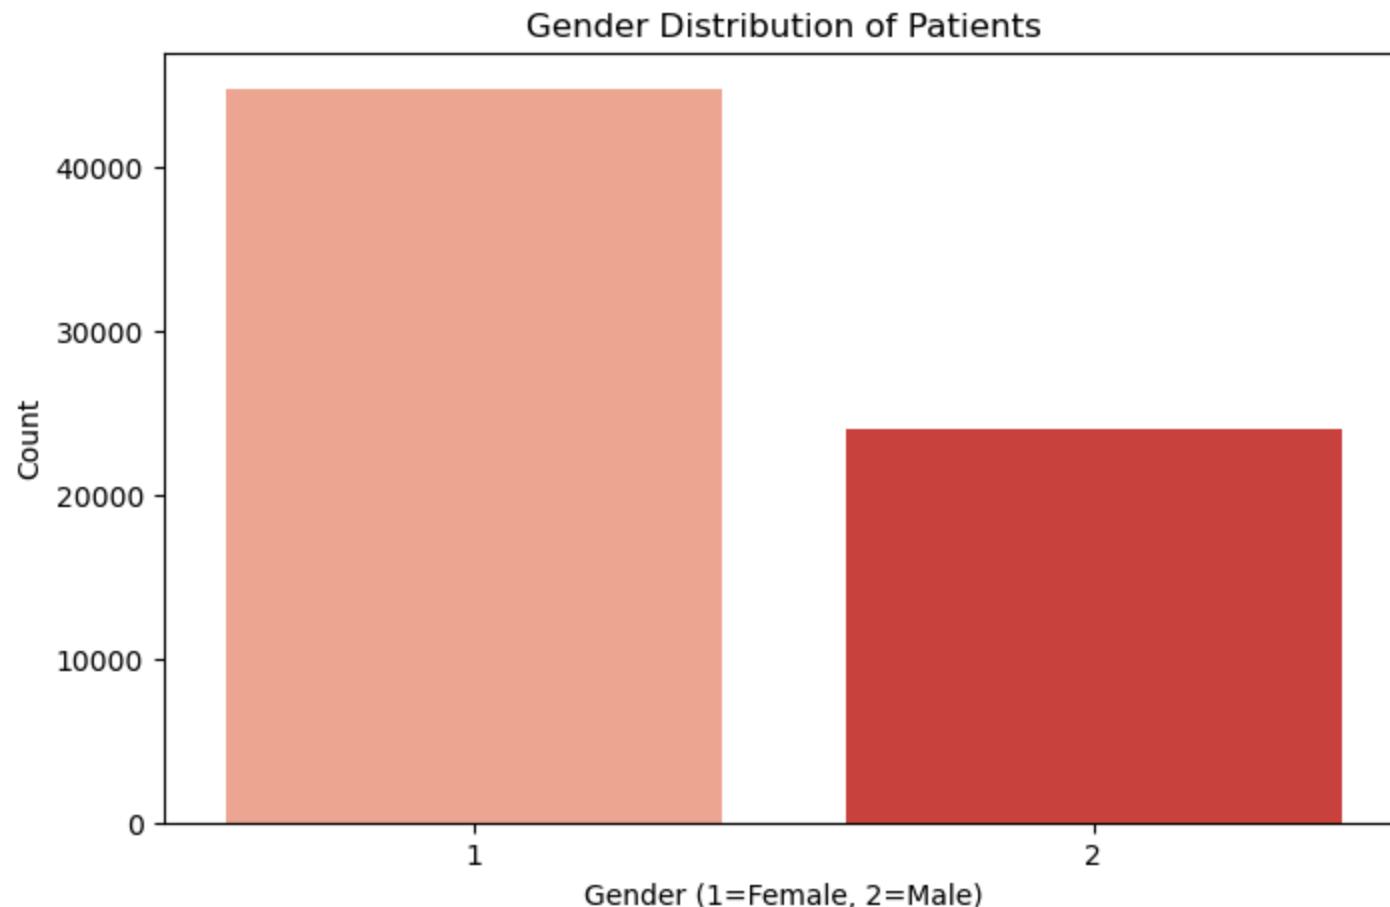
# 2. Demographics (Age, Gender)

```
In [12]: #A: Age Distribution  
# Plot age distribution  
plt.figure(figsize=(8,5))  
sns.countplot(x = 'age', data=df, color='Green')  
plt.title('Age Distribution of Patients')  
plt.xlabel('Age (years)')  
plt.ylabel('Count')  
plt.show()
```



```
In [13]: #B: Gender  
  
# Plot age distribution  
plt.figure(figsize=(8,5))
```

```
sns.countplot(x = 'gender', data=df, palette='Reds')
plt.title('Gender Distribution of Patients')
plt.xlabel('Gender (1=Female, 2=Male)')
plt.ylabel('Count')
plt.show()
```



### 3. Anthropometrics (Height, Weight, BMI)

In [14]:

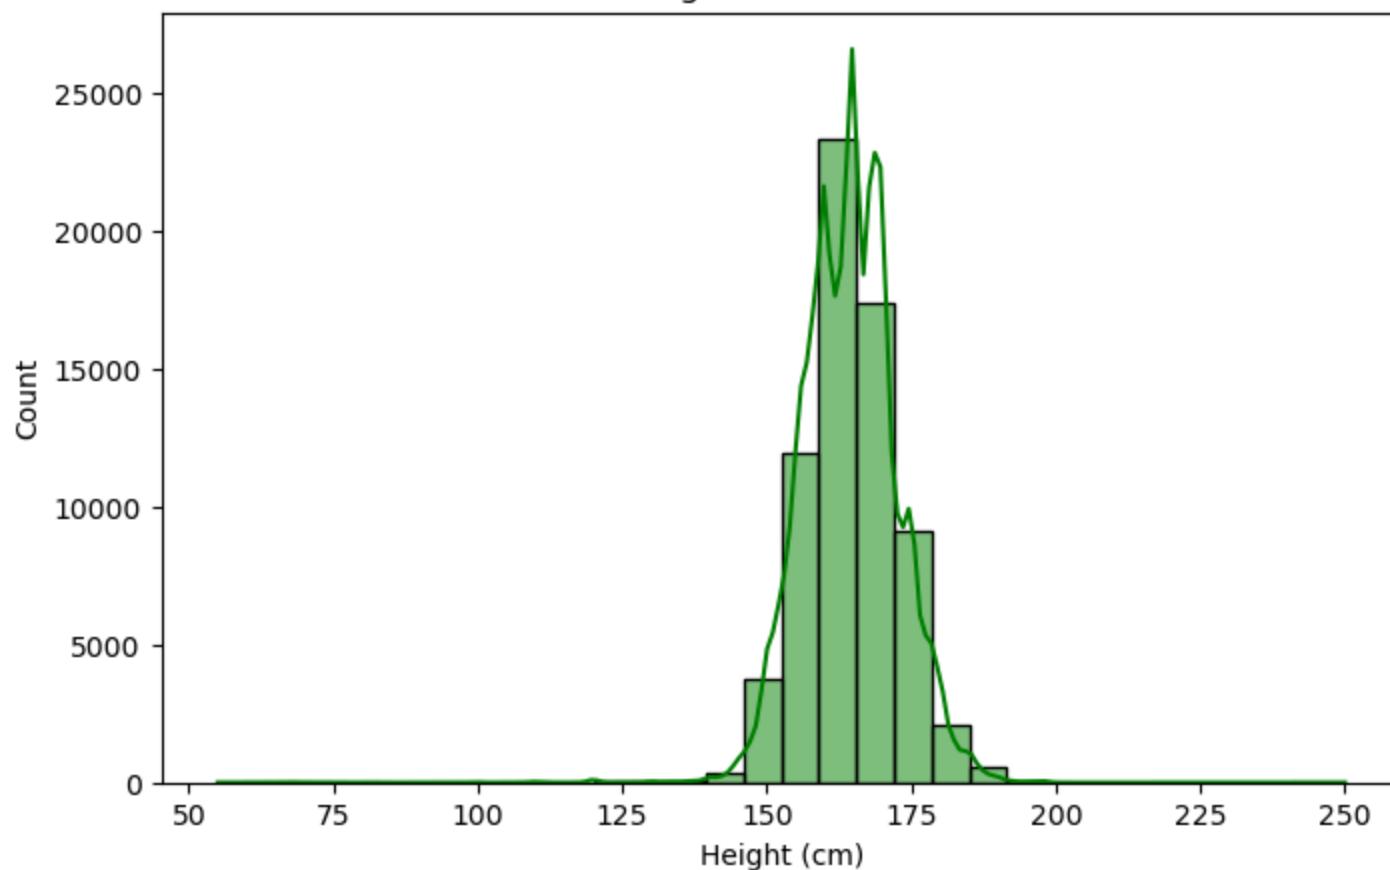
```
#Height Distribution
plt.figure(figsize=(8,5))
sns.histplot(df['height'], bins=30, kde=True, color='green')
```

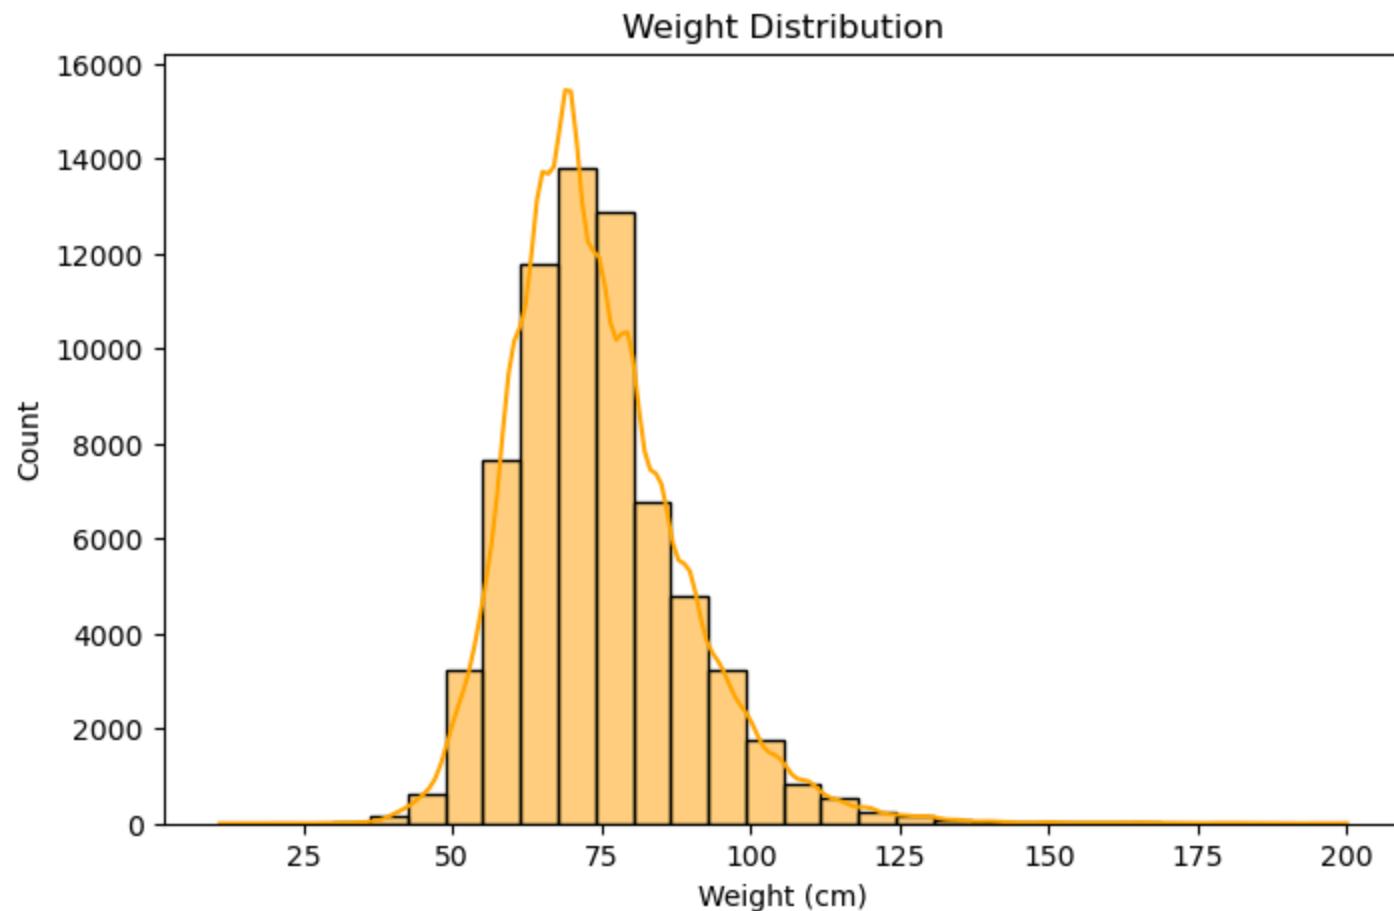
```
plt.title("Height Distribution")
plt.xlabel("Height (cm)")
plt.ylabel("Count")
plt.show()

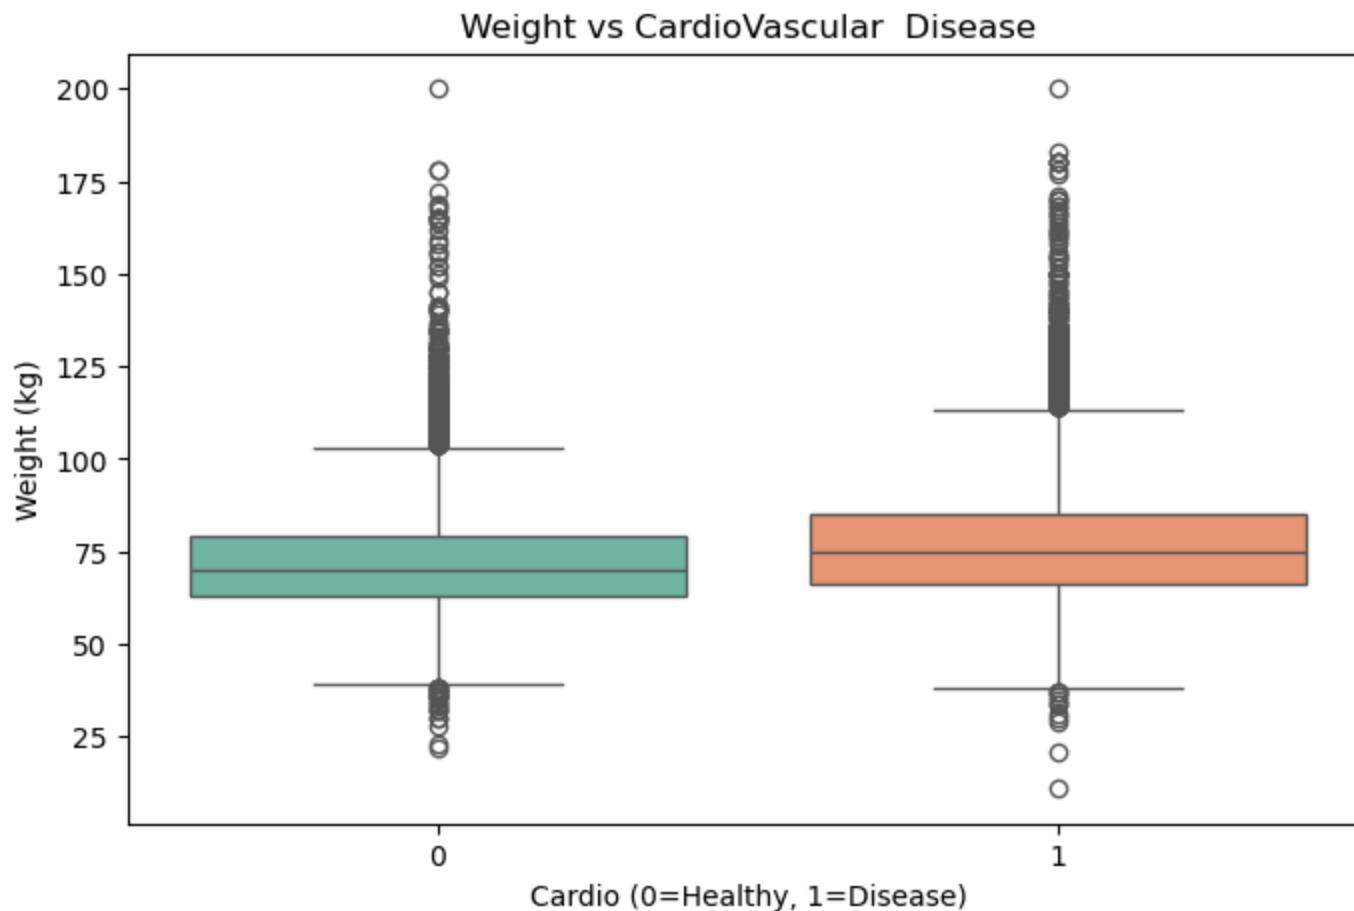
#Weight Distribution
plt.figure(figsize=(8,5))
sns.histplot(df['weight'], bins=30, kde=True, color='orange')
plt.title("Weight Distribution")
plt.xlabel("Weight (cm)")
plt.ylabel("Count")
plt.show()

#Compare Weight vs Cardio
plt.figure(figsize=(8,5))
sns.boxplot(x='cardio', y="weight", data=df, palette='Set2')
plt.title("Weight vs CardioVascular Disease")
plt.xlabel("Cardio (0=Healthy, 1=Disease)")
plt.ylabel("Weight (kg)")
plt.show()
```

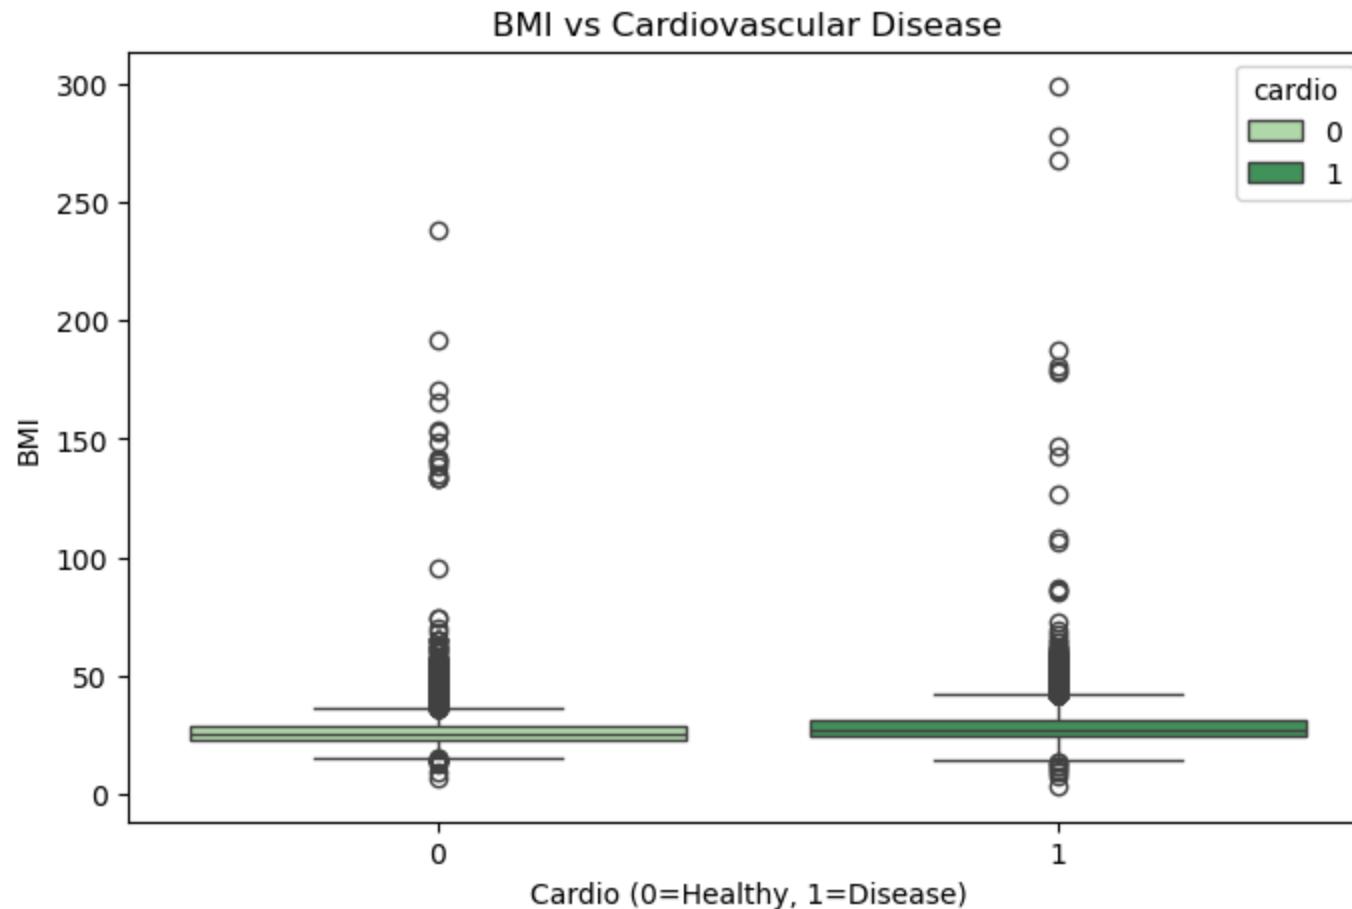
## Height Distribution







```
In [15]: # Compare BMI by disease status
plt.figure(figsize=(8,5))
sns.boxplot(data= df, x='cardio', y='bmi', hue='cardio', palette='Greens')
plt.title("BMI vs Cardiovascular Disease")
plt.xlabel("Cardio (0=Healthy, 1=Disease)")
plt.ylabel("BMI")
plt.show()
```

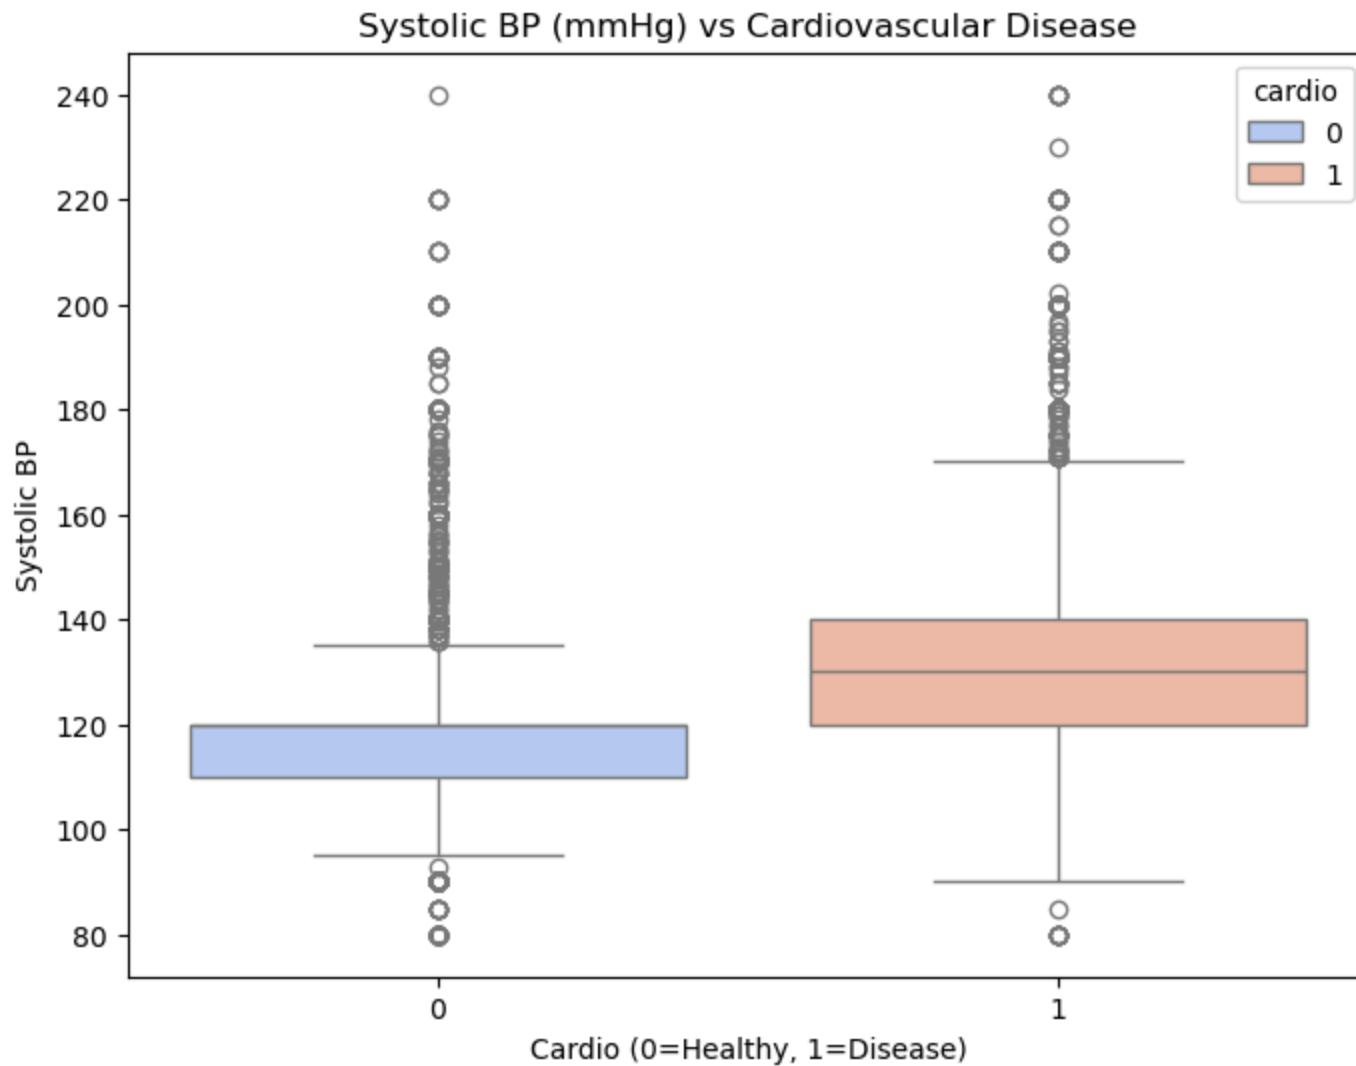


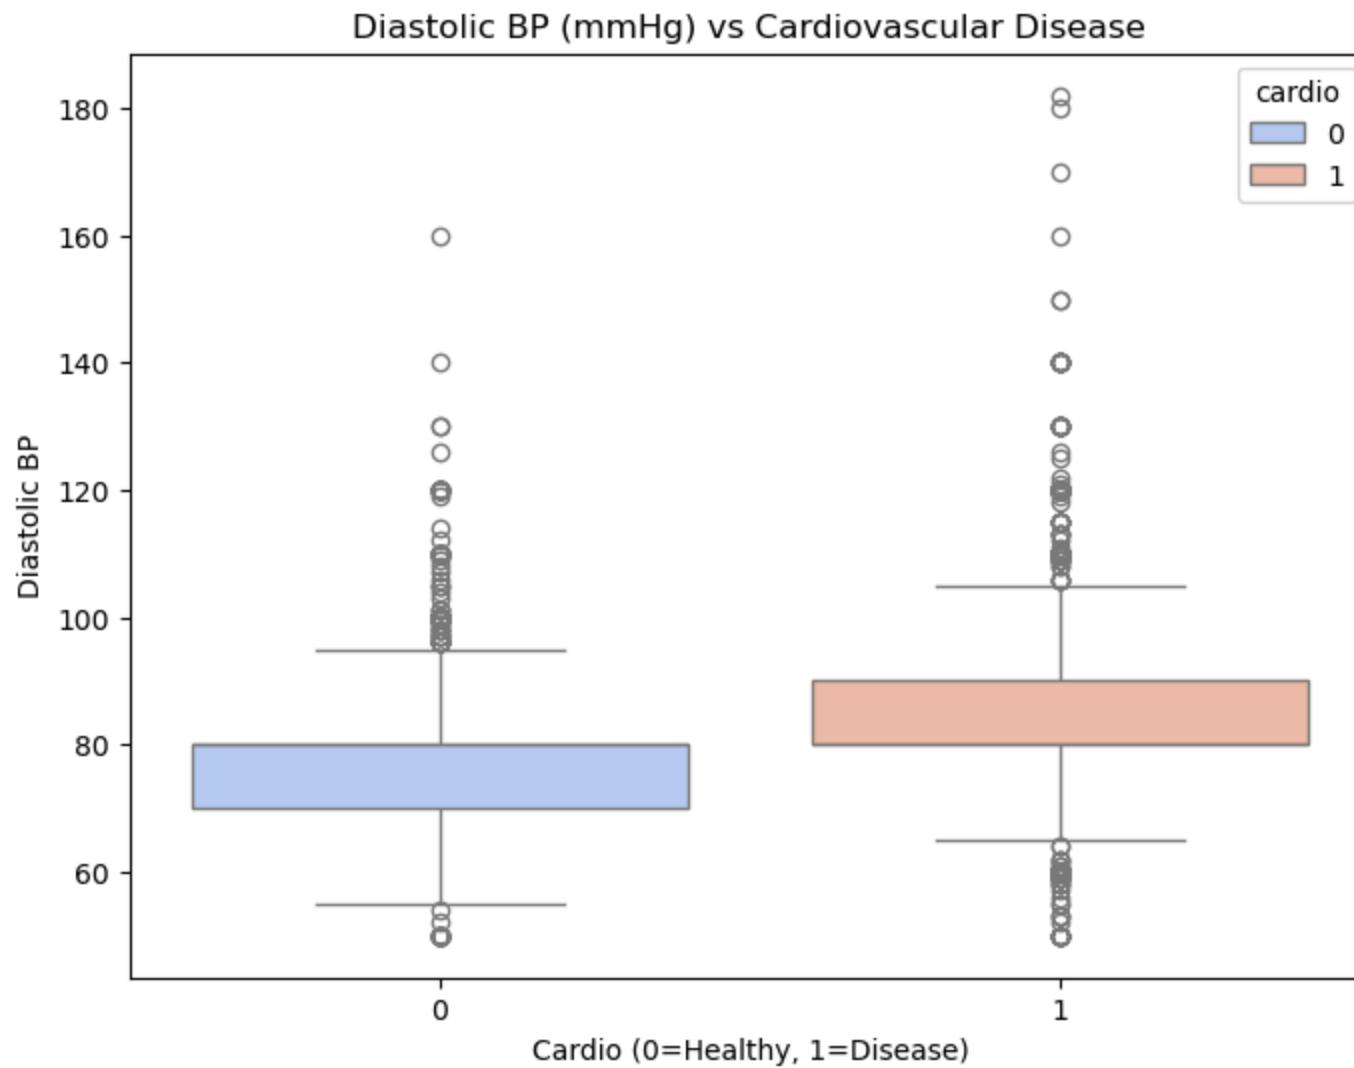
## 4. Medical Indicators (Blood Pressure, Cholesterol, Glucose)

```
In [16]: # Compared Blood Pressure by Disease
plt.figure(figsize=(8,6))
sns.boxplot(data= df, x='cardio', y='ap_hi', hue='cardio', palette='coolwarm')
plt.title("Systolic BP (mmHg) vs Cardiovascular Disease")
plt.xlabel("Cardio (0=Healthy, 1=Disease)")
plt.ylabel("Systolic BP")
plt.show()

plt.figure(figsize=(8,6))
```

```
sns.boxplot(data= df, x='cardio', y='ap_lo', hue='cardio', palette='coolwarm')
plt.title("Diastolic BP (mmHg) vs Cardiovascular Disease")
plt.xlabel("Cardio (0=Healthy, 1=Disease)")
plt.ylabel("Diastolic BP")
plt.show()
```





## 5. Lifestyle Factors (Smoke, Alcohol, Activity)

```
In [17]: #Smoking vs Disease  
#Create a temporary dataframe for plotting  
df_plot = df.copy()  
df_plot['smoke'] = df_plot['smoke'].map({0: "No", 1: "Yes"})  
df_plot['cardio'] = df_plot['cardio'].map({0: "Healthy", 1: "Disease"})
```

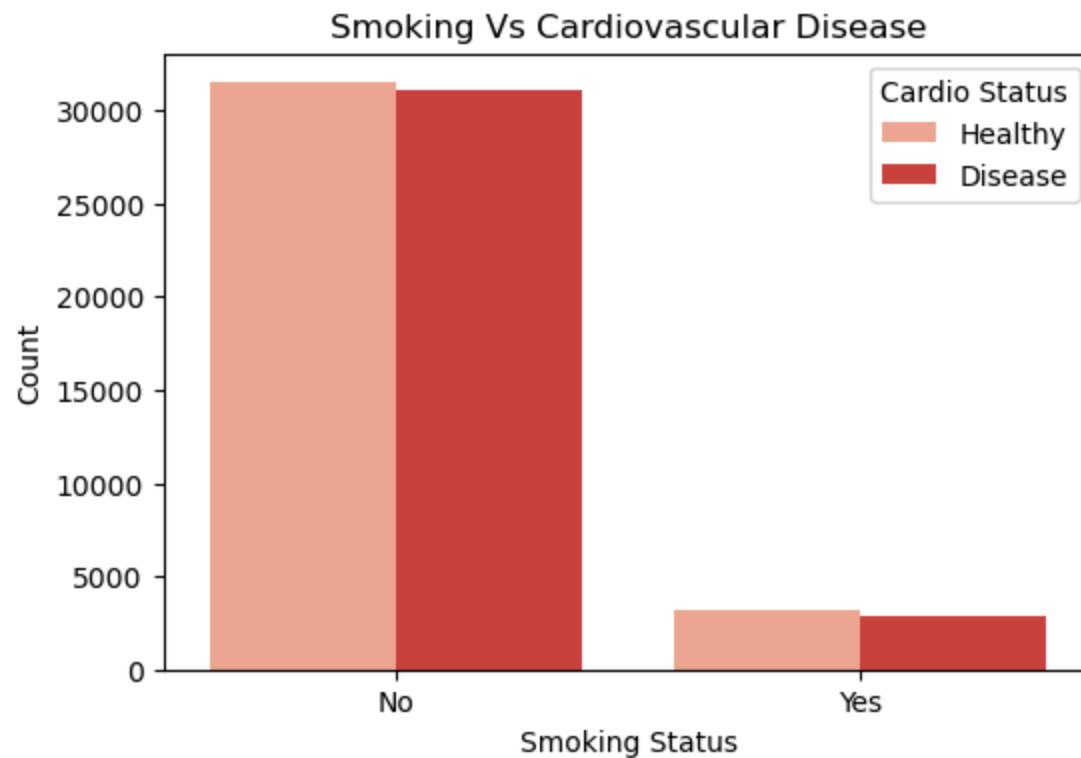
```
#Smoking vs Disease
plt.figure(figsize=(6,4))
sns.countplot(x='smoke', hue='cardio', data=df_plot, palette='Reds')
plt.title("Smoking Vs Cardiovascular Disease")
plt.xlabel("Smoking Status")
plt.ylabel("Count")
plt.legend(title="Cardio Status")
plt.show()

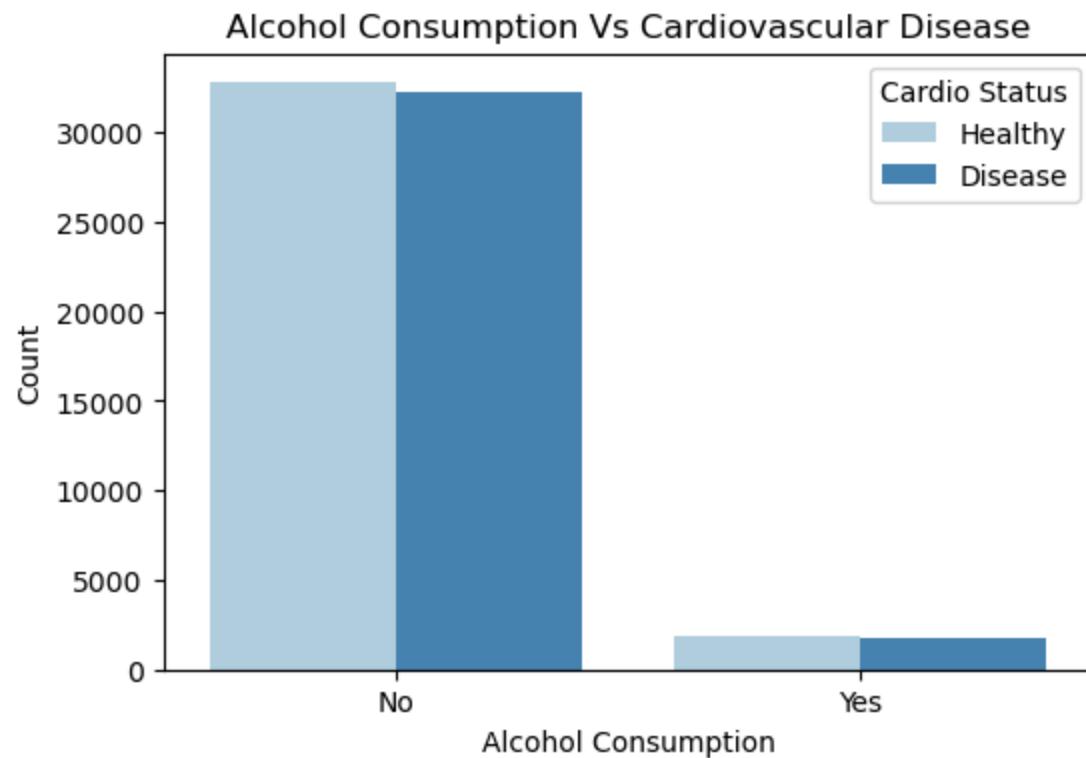
#Alcohol vs Disease
#Create a temporary dataframe for plotting
df_plot = df.copy()
df_plot['alco'] = df_plot['alco'].map({0: "No", 1: "Yes"})
df_plot['cardio'] = df_plot['cardio'].map({0: "Healthy", 1: "Disease"})

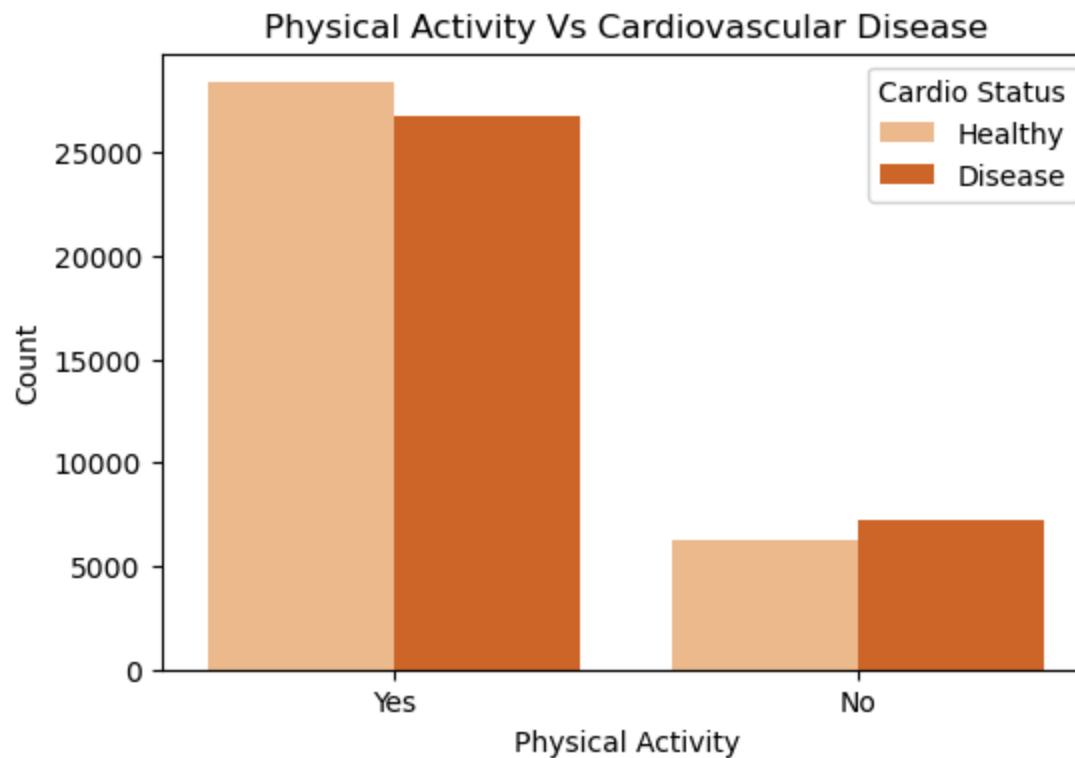
#Alcohol vs Disease
plt.figure(figsize=(6,4))
sns.countplot(x='alco', hue='cardio', data=df_plot, palette='Blues')
plt.title("Alcohol Consumption Vs Cardiovascular Disease")
plt.xlabel("Alcohol Consumption")
plt.ylabel("Count")
plt.legend(title="Cardio Status")
plt.show()

#Physical Activity vs Disease
df_plot['active'] = df_plot['active'].map({0: "No", 1: "Yes"})

#Physical Activity vs Disease
plt.figure(figsize=(6,4))
sns.countplot(x='active', hue='cardio', data=df_plot, palette='Oranges')
plt.title("Physical Activity Vs Cardiovascular Disease")
plt.xlabel("Physical Activity")
plt.ylabel("Count")
plt.legend(title="Cardio Status")
plt.show()
```







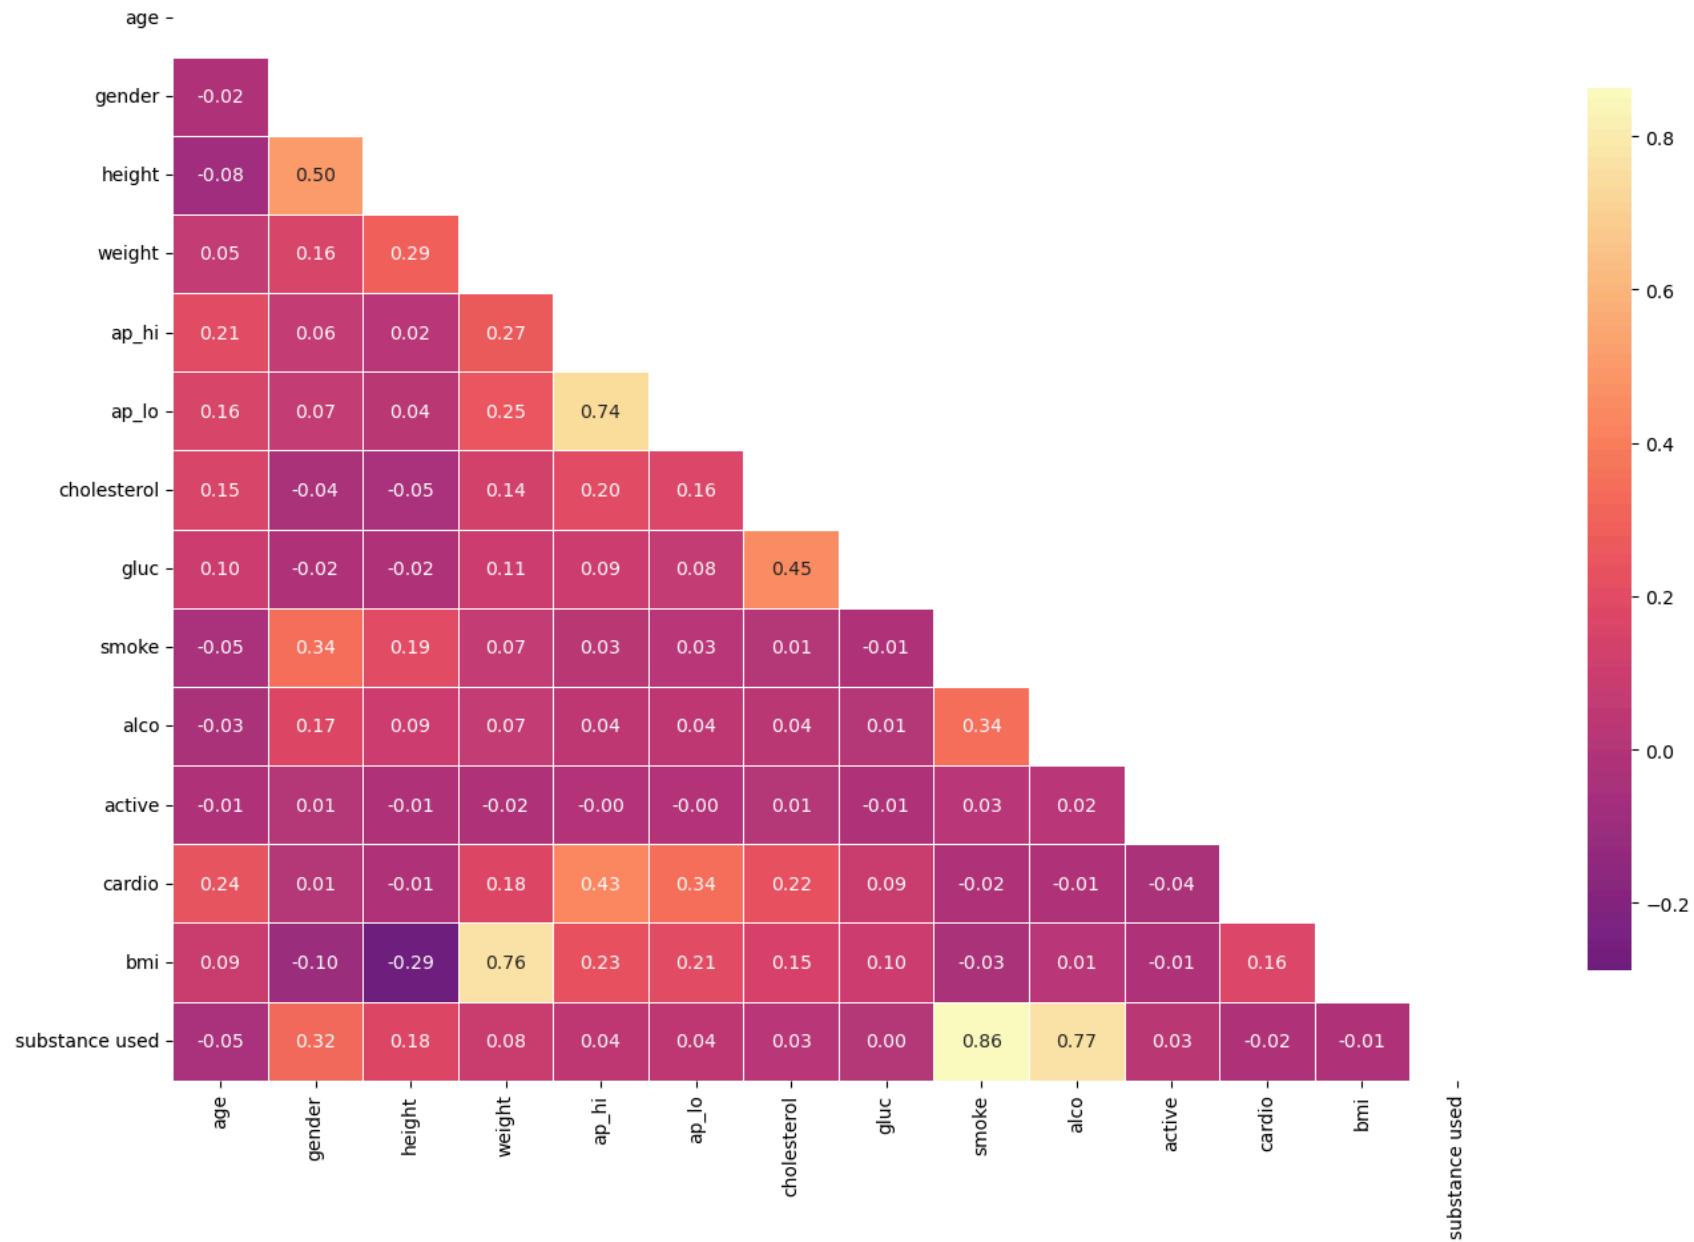
## 6. Correlation

```
In [18]: # --- Correlation Heatmap -
plt.figure(figsize=(14,10))
# Select only numeric columns for correlation
corr = df.select_dtypes(include=['number']).corr()
# Focus on stronger correlations only
mask = np.triu(np.ones_like(corr, dtype=bool)) # mask upper triangle
sns.heatmap(
    corr,
    mask=mask,
    cmap="magma",
    center=0,
    annot=True,
    fmt=".2f",
    linewidths= 0.5,
```

```
    cbar_kws={"shrink": 0.8}
)
plt.title("Correlation HeatMap for the Numeric Features of Heart Disease Attack", fontsize=14)
plt.tight_layout()
plt.show()

#Show correlation of features with target 'cardio'
target_corr = corr['cardio'].sort_values(ascending=False)
print("Correlation of Features with Target (cardio):\n")
print(target_corr)
```

## Correlation HeatMap for the Numeric Features of Heart Disease Attack

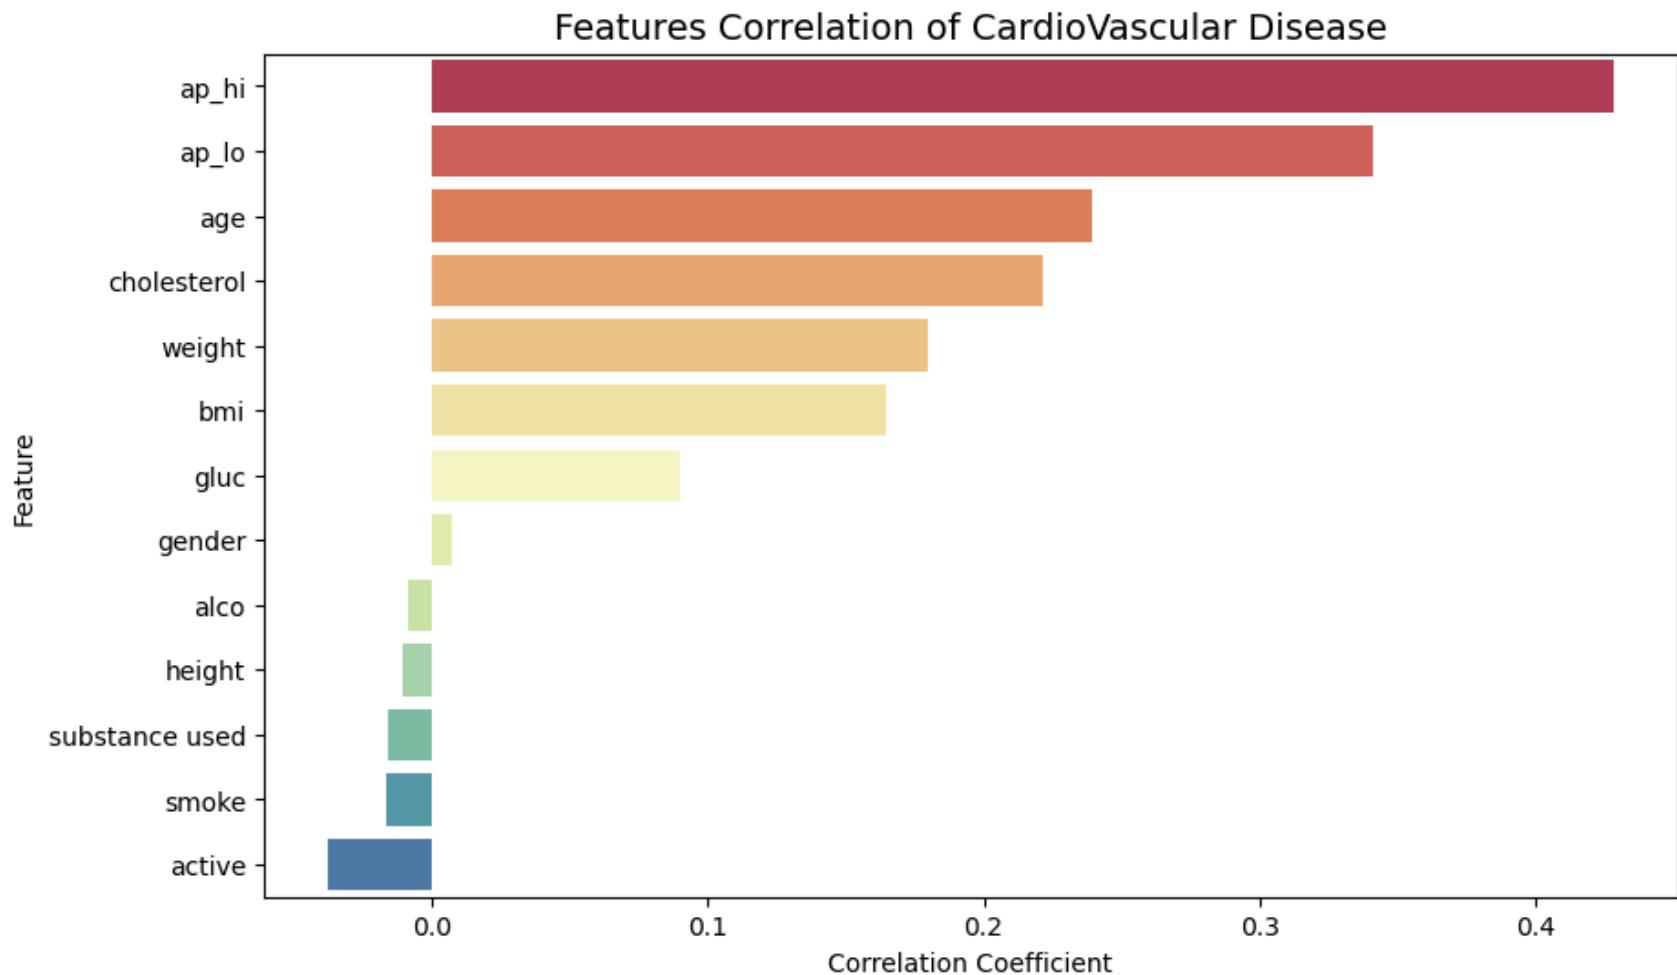


Correlation of Features with Target (cardio):

```
cardio      1.000000
ap_hi       0.427838
ap_lo       0.341166
age         0.239069
cholesterol 0.221410
weight      0.179623
bmi         0.164265
gluc        0.089658
gender      0.007127
alco        -0.008639
height      -0.010773
substance used -0.015656
smoke       -0.016207
active      -0.037558
Name: cardio, dtype: float64
```

## 7. Feature Importance

```
In [19]: #Sorted Correlated with target
numeric_df = df.select_dtypes(include=['number'])
target_corr = (numeric_df.corr()["cardio"]
               .drop('cardio')
               .sort_values(ascending=False)
               )
# --- Visualize correlation strength --
plt.figure(figsize=(10, 6))
sns.barplot(x=target_corr, y=target_corr.index, palette='Spectral')
plt.title("Features Correlation of CardioVascular Disease", fontsize=14)
plt.xlabel("Correlation Coefficient")
plt.ylabel('Feature')
plt.show()
```



## SPLITTING THE DATASET INTO TRAINING AND TESTING DATASET

```
In [20]: #Separate features (X) and target (y)

X = df.drop(columns='cardio')
y = df['cardio']

#Splitting the dataset into training (80%) and testing (20%)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42, stratify=y)
```

## IMPORTING MODEL

```
In [21]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
```

```
In [22]: lr = LogisticRegression(max_iter=500)
dt = DecisionTreeClassifier(random_state=42)
rf = RandomForestClassifier(random_state=42)
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
knn = KNeighborsClassifier()
gnb = GaussianNB()

models = {'Logistic Regression': lr, 'Decision Tree': dt, 'Random Forest': rf, "XGBoost": xgb, "KNeighborsClassifier": knn, 'Naive-Bayes': gnb}

for name, model in models.items():
    model.fit(X_train, y_train)
    print(f'{name} trained.')
```

```
Logistic Regression trained.
Decision Tree trained.
Random Forest trained.
XGBoost trained.
KNeighborsClassifier trained.
Naive-Bayes trained.
```

## Model Evaluation Use Metrics: Accuracy, Precision, Recall, F1\_Score and AUC\_ROC

```
In [23]: for name, model in models.items():
    y_pred = model.predict(X_test)
    print(f'{name}:')
    print(f'Accuracy: {accuracy_score(y_test, y_pred):.3f}')
    print(f'Precision: {precision_score(y_test, y_pred):.3f}')
    print(f'Recall: {recall_score(y_test, y_pred):.3f}')
    print(f'F1 Score: {f1_score(y_test, y_pred):.3f}')
    print(f'AUC-ROC: {roc_auc_score(y_test, model.predict_proba(X_test)[:,1]):.3f}\n')
```

Logistic Regression:

Accuracy: 0.724  
Precision: 0.748  
Recall: 0.668  
F1 Score: 0.706  
AUC-ROC: 0.785

Decision Tree:

Accuracy: 0.635  
Precision: 0.634  
Recall: 0.621  
F1 Score: 0.628  
AUC-ROC: 0.637

Random Forest:

Accuracy: 0.703  
Precision: 0.704  
Recall: 0.689  
F1 Score: 0.696  
AUC-ROC: 0.763

XGBoost:

Accuracy: 0.731  
Precision: 0.748  
Recall: 0.688  
F1 Score: 0.717  
AUC-ROC: 0.793

KNeighborsClassifier:

Accuracy: 0.688  
Precision: 0.691  
Recall: 0.668  
F1 Score: 0.679  
AUC-ROC: 0.735

Naive-Bayes:

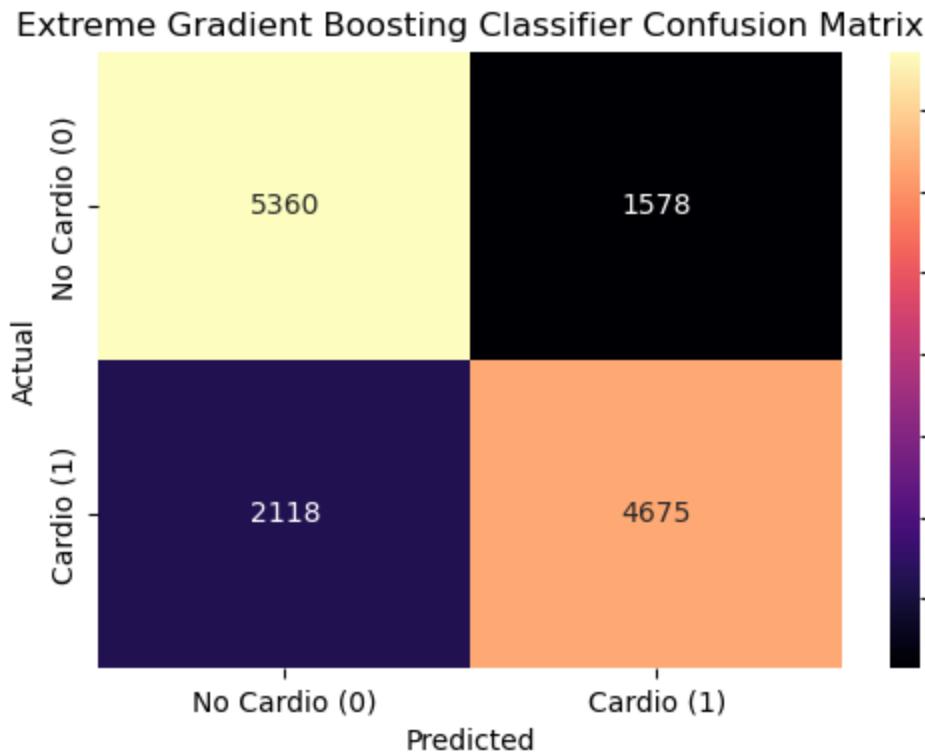
Accuracy: 0.707  
Precision: 0.753  
Recall: 0.606  
F1 Score: 0.671  
AUC-ROC: 0.772

# CONFUSION MATRIX

```
In [24]: model_xgb = xgb

#Confusion Matrix
y_pred_xgb = model_xgb.predict(X_test)
cm = confusion_matrix(y_test, y_pred_xgb)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='magma', xticklabels=['No Cardio (0)', 'Cardio (1)'],
            yticklabels=['No Cardio (0)', 'Cardio (1)'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Extreme Gradient Boosting Classifier Confusion Matrix')
plt.show()
```



1. True Positives (4675): These are individuals who were correctly identified as having cardiovascular disease. The organization can take appropriate steps like offering medical attention, lifestyle interventions, or further screening.
2. False Positives (1578): These individuals were incorrectly flagged as having the condition when they are actually healthy. This could lead to unnecessary follow-ups or concern but ensures caution.
3. False Negative(2118): These individuals actually have the condition, but the model failed to detect it. This is the most concerning group, as they may go untreated. Reducing this number is essential to prevent serious health outcomes.
4. True Negative (5360): These are healthy individuals correctly identified as such. No action is needed, helping focus resources where they are truly needed.

## APPLY STREAMLIT

```
In [25]: import joblib  
#save the trained model  
joblib.dump(lr, 'best_model_lr.pkl')  
print('Logistic Regression model saved as best_model_lr.pkl')
```

Logistic Regression model saved as best\_model\_lr.pkl

```
In [26]: import joblib  
#save the trained model  
joblib.dump(xgb, 'best_model_xgb.pkl')  
print('Extreme Gradient Boosting Classifier model saved as best_model_xgb.pkl')
```

Extreme Gradient Boosting Classifier model saved as best\_model\_xgb.pkl